

# Hand Gesture Recognition

Team members: Paul Woodward, Sahiti Tadepalli, Hien Nguyen, Annie Liang, Xinyi Zhang

-----Machine Learning , Winter 2018

# Outline

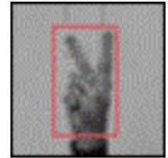
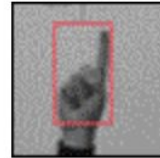
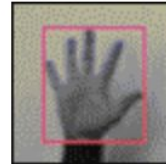
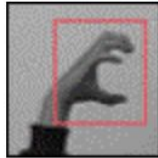
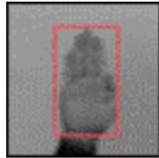
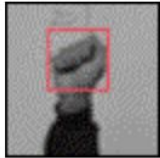
- Project intro
- Design Overview
- Model Detail
- Results
- Morbidity and Mortality
  - tuning attempts

---

# Project Intro:

## Sebastien Marcel Static Hand Posture Database:

6 hand postures (a, b, c, point, five, v), about 10 persons.



Count of the number of images in each section:

Found 4872 images belonging to 6 classes (training set)

Found 659 images belonging to 6 classes (test set)

Original Image sizes:

{'76,66,3': 1573, '82,70,3': 471, '84,72,3': 147, '90,78,3': 495, '288,384,3': 153, '155,155,3': 86, '320,240,3': 82, '88,76,3': 461, '80,68,3': 499, '66,56,3': 116, '78,66,3': 199, '74,64,3': 169, '66,76,3': 110, '68,80,3': 59, '76,88,3': 55, '70,82,3': 51, '72,84,3': 24, '64,80,3': 6, '66,80,3': 11, '78,90,3': 79, '187,194,3': 5}

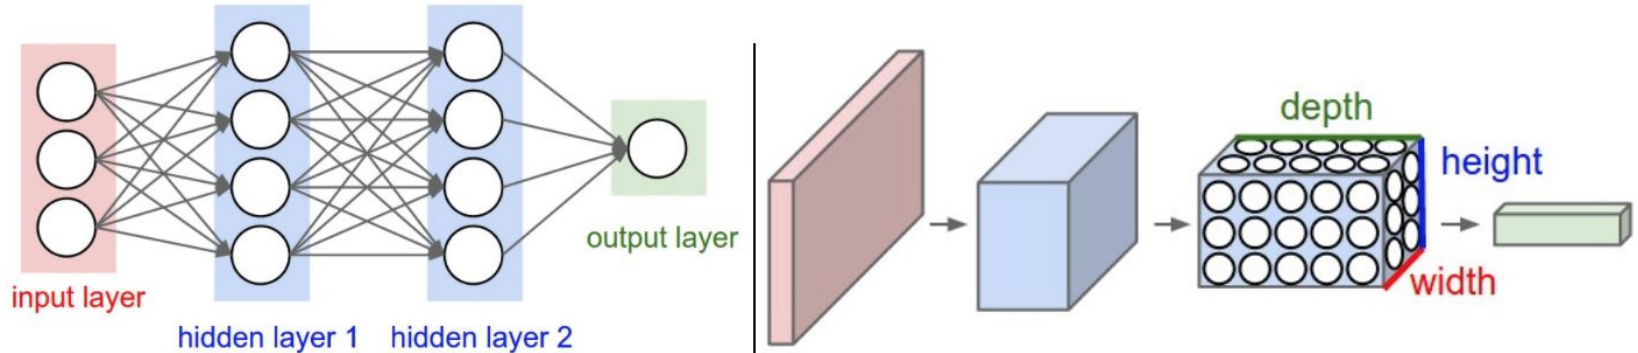
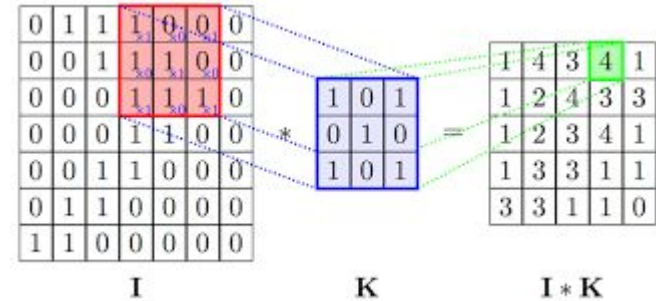
Pre-processed: 64 x 64 x 1 ( also the input pixel values)

Filter/Receptive field: 3 x 3 x 1

CONV layer (hidden layer): slide over the image spatially,  
computing dot products

- stride( 2,2)

$$\mathbf{W}^T \mathbf{x} + b$$

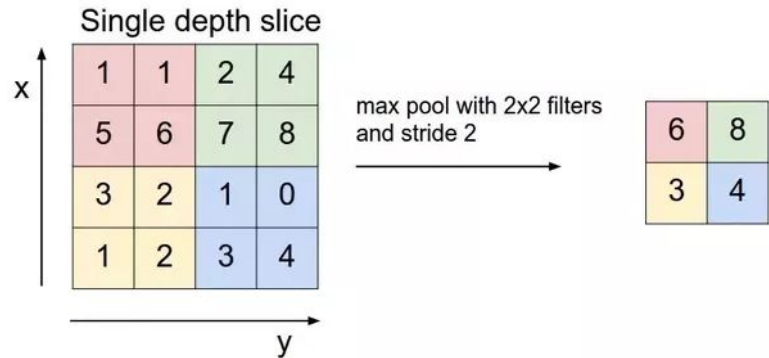
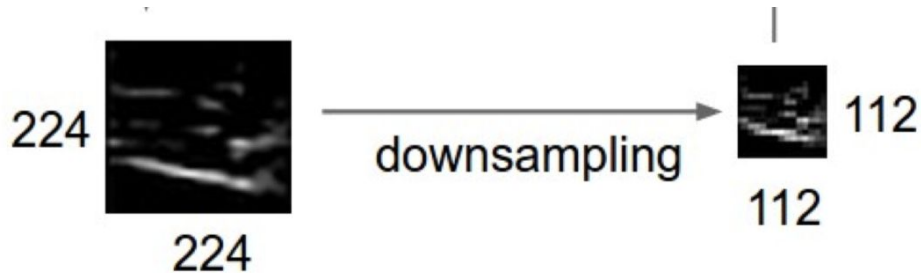


Output size is computed using  $(N-F)/S + 1$ .

Here, we have 64x64 input and 3x3 filter with 2 stride gives us 32 x32.

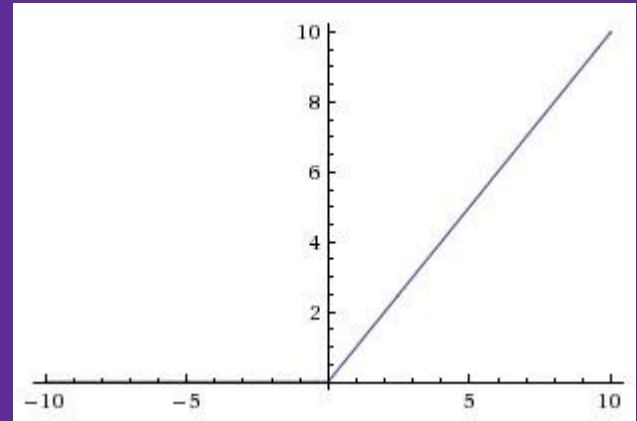
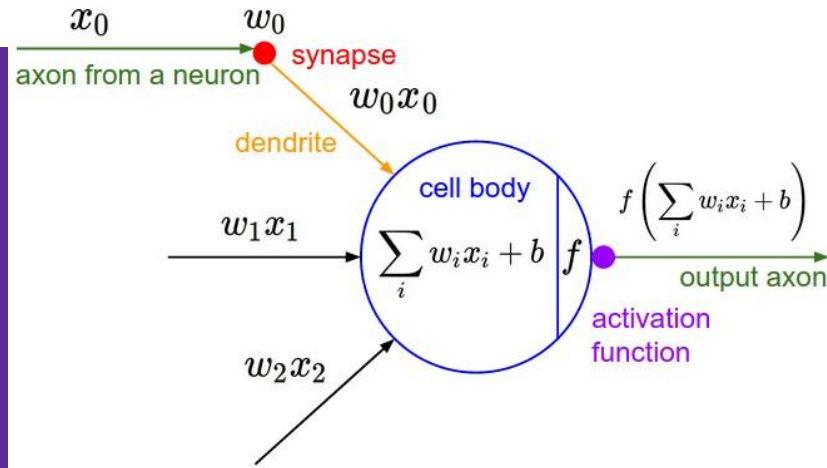
This “new image” is the input for Maxpooling function that takes only max value per the grid.

- dimension reduction. It only keeps the most intense pixels but also retain image information



# Activation Function

rectified linear unit (ReLU)



# Model summary:

```
model.summary()
```

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 62, 62, 32)	320
max_pooling2d_1 (MaxPooling2D)	(None, 31, 31, 32)	0
activation_1 (Activation)	(None, 31, 31, 32)	0
batch_normalization_1 (Batch Normalization)	(None, 31, 31, 32)	128
conv2d_2 (Conv2D)	(None, 29, 29, 32)	9248
activation_2 (Activation)	(None, 29, 29, 32)	0
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 32)	0
flatten_1 (Flatten)	(None, 6272)	0
dense_1 (Dense)	(None, 128)	802944
dropout_1 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 6)	774
=====		
Total params: 813,414		
Trainable params: 813,350		
Non-trainable params: 64		

# Model Parameters

Filter: size 3x3

Activation function:

- Relu: for Convolutional Layers, MaxPooling Layer and Dense except Fully Connected Layer.
- Softmax: for the fully Connected layer in Multiple Label classification. Softmax was built for the purpose of probabilistic classification.

Optimizer:

- SGD and learning rate start at 5e-3 and adjust over time using LearningRateSchedule function built in Keras



# Results:



# Model accuracy:

```
2 model.fit_generator(training_set, steps_per_epoch = TRAIN_IMG_CNT / batch_size, epochs = NUM_EPOCHS,  
3                     callbacks=callbacks,  
4                     validation_data = test_set, validation_steps = VALIDATION_CNT / batch_size)  
5 model.save_weights('1st_run.h5')
```

Epoch 91/100

305/304 [=====] - 10s 33ms/step - loss: 0.0743 - acc: 0.9736 - val\_loss: 1.3238 - val\_acc: 0.7754

Epoch 92/100

305/304 [=====] - 10s 33ms/step - loss: 0.0737 - acc: 0.9713 - val\_loss: 1.4556 - val\_acc: 0.7557

Epoch 93/100

305/304 [=====] - 10s 33ms/step - loss: 0.0789 - acc: 0.9721 - val\_loss: 1.4525 - val\_acc: 0.7451

Epoch 94/100

305/304 [=====] - 10s 32ms/step - loss: 0.0664 - acc: 0.9742 - val\_loss: 1.5095 - val\_acc: 0.7329

Epoch 95/100

305/304 [=====] - 10s 33ms/step - loss: 0.0744 - acc: 0.9721 - val\_loss: 1.4269 - val\_acc: 0.7542

Epoch 96/100

305/304 [=====] - 11s 35ms/step - loss: 0.0697 - acc: 0.9734 - val\_loss: 1.6327 - val\_acc: 0.7238

Epoch 97/100

305/304 [=====] - 10s 33ms/step - loss: 0.0710 - acc: 0.9746 - val\_loss: 1.6378 - val\_acc: 0.7269

Epoch 98/100

305/304 [=====] - 10s 34ms/step - loss: 0.0698 - acc: 0.9748 - val\_loss: 1.5895 - val\_acc: 0.7329

Epoch 99/100

305/304 [=====] - 10s 33ms/step - loss: 0.0623 - acc: 0.9768 - val\_loss: 1.6094 - val\_acc: 0.7344

Epoch 100/100

305/304 [=====] - 10s 34ms/step - loss: 0.0727 - acc: 0.9721 - val\_loss: 1.5964 - val\_acc: 0.7375

# Using the trained model to make some predictions:

In [29]:

```
1
2 for i in range(11,13):
3     for ch in ['A','B', 'C']:
4         filename = "Marcel-Test/" + ch + "/uniform/" + ch + (
5             "-uniform" + str(i) + ".ppm")
6         y = predict(filename)
7         print(ch, y)
```

A [0]

B [1]

C [2]

A [3]

B [1]

C [2]

# Morbidity and Mortality



64 32 128 6 - 70% kern=5x5  
256 128 64 6 - slow  
256 64 6 - slow, bad  
128 64 6 - bad  
32 128 6 - 57%  
32 128 6 kern=3x3 - 45%  
32 128 6 kern=7x7 - 50%  
32 128 6 batch size = 32, kern = 5, opt=adam - 50%  
64 128 6 batch 16 slow  
16 128 6 - 55  
del relu from first cov2d 50  
32 32 128 6 kern=3 70 del white  
32 32 256 6 - 65  
32 32 192 6 - 70+  
32 32 192 6 with poly\_decay - 74  
32 32 192 6 with more-train - meh  
32 32 256 6 with marcel-train - 71  
32 32 160 - 68  
32 32 32! 160 6 - 61  
32 16 16 64 - no|  
16 16 16 128 - no  
64 16 16 128 - no  
32 32 32 190 - no  
32 32 192 - ok, 70  
32 32 128 - good  
24 24 190 - 70  
24 24 190 no noise : overfit  
24 24 190, 20% noise + 2 degrees: nope  
10% and 0 degrees... - 73  
32 32 192 2 degrees... 70  
32 32 192 1 ... 73  
32 32 128 1 ... 73



File Edit Format View Help

failed:

image processing to edges, position -- never started

tensor flow -- it never ran

aws -- no gpu for you

google cloud -- anomaly detected!

no noise -- overfitting

more noise -- no accuracy

running noise once -- overfitting

large epoch -- too slow

lots of nodes -- overfitting

even more layers -- no effect

???:

image size / scaling -- losing data going to 64x64?

monochrome vs color -- early tests favored grayscale

convolution & pooling on hidden layers (seems like image processing)

good:

installing Tensorflow GPU -- fast

keras -- intelligent defaults, got us going

**Thank you**